

Dissertation Prospectus

A Full-Spectrum Dependently typed language for testing with dynamic equality

Mark Lemay

Department of Computer Science, Boston University

January 28, 2021

Abstract

Dependent type systems are a powerful tool to eliminate bugs from programs. Unlike other systems of formal methods, dependent types systems can re-use the methodology and syntax that functional programmers are already familiar with for the construction of formal proofs. This insight has lead to several Full-Spectrum languages that try and present programmers with a consistent view of proofs and programs. However these languages still have substantial usability issues: missing features like general recursion, confusingly conservative equality, and no straight forward way to test specifications that have not yet been proven.

These issues are not superficial, but are tied to some of the conventional assumptions of dependent type theory. Logical soundness is essential when dependent types are considered as a logic, but type soundness is sufficient when used as a programming language. The philosophical commitment of equality is easy to hand-wave away informally, but it makes many programs unrunnable. Users often experience these assumptions as inexplicable static errors. In practice these foundational assumptions, make exploratory programming difficult.

A possible solution is to design a dependently typed language with programmatic features that supports type sound, but logically unsound execution. Programmers can trade poor static errors, for precise counter examples that are made available at runtime. Progress towards that solution has been by developing a base language with programming oriented features, characterizing some correctness criteria, and implementing a subset of that language that matches those criteria.

Contents

1 Introduction

2

2	A Dependently Typed Base Language	4
2.1	Prior work for the Base Language	5
3	A Language with Dynamic Type Equality	6
3.1	Status	8
3.2	Prior work	8
4	Future Work	9
4.1	Prototyping proofs and programs	9
4.1.1	Prior work	10
4.2	Testing dependent programs	11
4.2.1	Prior work	11
4.2.2	Symbolic Execution	11
4.2.3	Testing dependent types	12

1 Introduction

The Curry-Howard correspondence identifies functions with theorems, providing a mutually beneficial link between well explored areas of math and software engineering. This connection is most pronounced in dependent typed systems that provide a common language to write programs about proofs and proofs about programs. Specifically, dependent type systems allow types to depend on terms, supporting an extremely powerful system of specification, computation and proof evidence.

For instance, in a dependently typed language it is possible to prove the correctness of a sorting function

$$\text{sort} : (\text{input} : \text{List } \mathbb{N}) \rightarrow \Sigma ls : \text{List } \mathbb{N}. \text{IsSorted input } ls$$

by providing an appropriate term of that type. Unlike other systems of formal methods, the additional logical power does not require the programmer understand any additional syntax or semantics. From the programmer's perspective the function arrow and the implication arrow are the same. The proof *IsSorted* is no different then any other datatype like *List* or \mathbb{N} .

The promise of dependent types in a practical programming language has been the goal of research projects for

decades. There have been many formalization and prototypes that make different compromises in the design space. One of the most popular styles is so called “Full-spectrum” dependently types languages, these languages tend to have a minimalist approach: computation can appear anywhere in a term or type. Such a design purposely exposes the Curry-Howard correspondence, as opposed to trying to hide it as a technical foundation: a proof has the exact same syntax and behavior as a program. Even though this style makes writing efficient programs hard, and drastically complicates the ability to encode effects, it can be seen in some of the most popular dependently typed languages such as Agda and Idris.

Despite the potential, users often find these systems difficult to use. The common symptom of these issues can be seen in the confusing error messages these languages produce. For instance in Agda this reasonable looking program

$$\begin{aligned} \text{Vec} &: * \rightarrow \mathbb{N} \rightarrow * \\ \text{rep} &: (x : \mathbb{N}) \rightarrow \text{Vec } \mathbb{B} \, x \\ \text{head} &: (x : \mathbb{N}) \rightarrow \text{Vec } \mathbb{B} \, (1 + x) \rightarrow \mathbb{B} \\ &\not\vdash \lambda x. \text{head } x \, (\text{rep } (x + 1)) : \mathbb{N} \rightarrow \mathbb{B} \end{aligned}$$

will give the error “ $x + 1 \neq \text{suc } x$ of type \mathbb{N} when checking that the expression $\text{rep } (x + 1)$ has type $\text{Vec } \mathbb{B} \, (1 + x)$ ”. The error is confusing since it objects to an intended property of addition, and if addition were buggy no hints are given to fix the problem. While an expert in Type Theory can appreciate the subtitles of definitional equality, programmers would prefer an error message that gives a specific instance of x where $x + 1 \neq 1 + x$ or remain silent.

Strengthening the equality relation in dependently typed languages is used to motivate many research projects (to name a few [7, 27]). However, it is unlikely those impressive efforts are suitable for non-experts, since programmers expect the data types and functions they define to have the properties they were intended to have. None of these projects makes the underlying equality less subtle and complicated. There is not enough information available statically to verify every “obvious” equality for arbitrary user defined data types and functions.

Alternatively we could assume the equalities hold and discover a concrete inequality as a runtime error. There

is some evidence that specific examples like this can help clarify the error messages in OCaml[23] and there has been an effort to make refinement type error messages more concrete and other systems like Liquid Haskell[12]. This leads too a different workflow then traditional type systems, instead of verifying everything first and only then executing the program, execution and type checking can inform each other.

To make this new workflow possible” section 2 will describe a more traditional base language with some programmatic features. Section 3 describes dynamic dependent equality and the conditions that would make an implementation reasonable. Section 4 covers some of the other features a language in this style could support.

2 A Dependently Typed Base Language

Dynamic dependent equality is hard to study without a more conventional dependently typed language to serve as a reference. This base language contains the key features that should be supported: user defined dependent functions and dependent datatypes. The base language uses a “Full-Spectrum” style. The dynamic language from the next section can be elaborated from and compared to this concrete implementation.

The language is pure in the sense of Haskell, supporting only non-termination and unchecked errors as effects. Combing other effects with Full-Spectrum dependent types is substantially more difficult because effectful equality is hard to characterize for individual effects or effects in combination. Several attempts have been made to combine dependent types with more effects, [22] [2, 1][22] but there is still a lot of work to be done, so it is not considered in the base language.

Since this work emphasizes programming over theorem proving the language contains these logically dubious features:

- Unrestricted user defined dependent data types (no requirement of strict positivity)
- Unrestricted recursion (no required termination checking)
- Type-in-type (no predictive hierarchy of universes)

Any one of these features can result in logical unsoundness¹, but they facilitate mainstream functional programming. In spite of the logical unsoundness, the resulting language still has type soundness². This seems suitable for

¹Every type is inhabited by an infinite loop.

²No term with a reduct that applies an argument to a non-function in the empty context will type.

a programming language since logically sound proofs can still be defined and logical unsoundness can be discovered through traditional testing, or warned about in a non-blocking way. Importantly no desirable computation is prevented.

Though this language is not logically sound, it supports a partial correctness property for first order data types when run with Call-by-Value, for instance:

$$\vdash M : \sum x : \mathbb{N}. \text{IsEven } x$$

`fst M` may not terminate, but if it does, `fst M` will be an even \mathbb{N} . However, this property does not extend to functions

$$\vdash M : \sum x : \mathbb{N}. (y : \mathbb{N}) \rightarrow x \leq y$$

it is possible that `fst M` \equiv 7 if

$$M \equiv \langle 7, \lambda y. \text{loopForever} \rangle$$

The hope would be that the type is sufficient to communicate intent, in the same way unproductive non-termination is typeable in mainstream programming languages but still considered a bug.

2.1 Prior work for the Base Language

While many of these features have been explored in theory and implemented in practice, I am unaware of any development with exactly this formulation.

Unsound logical systems that are acceptable programming languages go back to at least to Church's lambda calculus which was originally intended to be a logical foundation for mathematics. Martin Lof proposed a system with Type-in-type that was shown logically unsound by Girard (as described in the introduction in [18]).

The first proof of type soundness for general recursive functions that I am aware of came from the Trellys Project [26], it contains many similar features, but the base language uses a simpler notion of equality and dependent data

resulting in an arguably simpler proof of type soundness. Further work in the Trellys Project[6, 5] used modalities to separate the terminating and non terminating fragments of the language, though the annotations burden seems too high in practice. In general the base language has been deeply informed by the Trellys Project[14][26][6, 5] [27] [25] and the Zombie Language³ it produced.

Many implementations support this combination of features without proofs of type soundness. Cayenne [4] was an early Haskell like language combined dependent types with and non-termination. Agda supports general recursion and Type-in-type with compiler flags, and can simulate some non-positive data types using coinduction. Idris supports similar “unsafe” features.

[13] claims a similar “partial correctness” criterion.

3 A Language with Dynamic Type Equality

A key issue with full-spectrum dependent type theories is the characterization of definitional equality. Since computation can appear at the type level, and types must be checked for equality, dependent type theories must define what computational equalities they intend to support. For instance, the base language follows the common choice of $\alpha\beta$ equivalence of terms. However this causes many, programs to not type-check:

$\lambda x.\text{head } x \ (\text{rep } (x + 1)) : \mathbb{N} \rightarrow \mathbb{B}$ Since $1 + x$ does not have the same definition as $x + 1$.

This is a widely recognized issue with Dependent Type Theories, However most attempts to improve the equality relation intend to preserve decidable type checking and/or logical soundness, so equality will never be complete⁴. Since dependently typed languages with the practical features outlined in base language are already incompatible with logical soundness and decidable type checking, these concerns no longer apply.

The base language can be extended to a cast language that supports the expectation of the original typing. Many programs that do not type in the base language can be elaborated into the cast language. The cast language has a weaker notion of type soundness such that

1. $\vdash_c e' : M'$ then

(a) $e' \downarrow v'$ and $\vdash_c v' : M'$

³<https://github.com/sweirich/trellys>

⁴I am also unaware of any suitable notion of complete extensional equality for dependent type theory though it is considered in [25] .

(b) or $e' \uparrow$

(c) or $e' \downarrow \text{blame}$

Type soundness is preserved, or inequality can be proven at a specific source location. In the example above $\lambda x.\text{head } x (\text{rep } (x + 1)) : \mathbb{N} \rightarrow \mathbb{B}$ will not emit any errors at compile time or runtime (though a static warning may be given).

If the example is changed to

$$\lambda x.\text{head } x (\text{rep } x) : \mathbb{N} \rightarrow \mathbb{B}$$

and the function is called at runtime, the blame tracking system will blame the exact static location that uses unequal types with a direct proof of inequality, allowing an error like “failed at application $(\text{head } x : \text{Vec } \underline{(1 + x)} \mathbb{B} \rightarrow \dots) (\text{rep } x : \mathbb{N})$ ” since when $x = 3$, $1 + x = 4 \neq 3 = x$, regardless of where function was called in the program the discrepancy was discovered. This improves on the naive solution of completely ignoring type annotations for execution, since without type soundness, the returned error may appear unrelated to the problematic type assumption (“3 applied to 7 but 7 is not a functions”).

Just as standard type theories allow many possible characterizations of equality that support logical soundness, there are many choices of runtime checking that support this notion of type soundness. The minimal choice that is likely too permissive in practice. In the example above, it would only flag an error when the function is applied to 0. Alternatively, I conjecture that runtime-checking that matches the partial correctness criteria above would be reasonably intuitive. Extending checks into non dependent function types also seems reasonable, and would allow simple types to be checked statically.

Taking inspiration from the “gradual guarantee” of gradual typing, there are several basic properties in addition to type soundness that this cast language hopes to fulfill:

1. $\vdash e : M$, $\text{elab } (M, *) = M'$, and $\text{elab } (e, M') = e'$ then $\vdash_c e' : M'$.
2. $\vdash_c e' : M'$ and $e' \downarrow \text{blame}$ then there is no $\vdash e : M$ such that $\text{elab } (M, *) = M'$, $\text{elab } (e, M') = e'$
3. $\vdash_c e' : *$ and $\text{elab } (e, *) = e'$ then

(a) if $e' \downarrow *$ then $e \downarrow *$

(b) if $e' \downarrow (x : M') \rightarrow N'$ then $e \downarrow (x : M) \rightarrow N$

(c) if $e' \downarrow TCon\Delta'$ then $e \downarrow TCon\Delta$

The first condition states that every typed term in the base language can be embedded in the cast language. The second condition shows that errors are not spurious. The third condition shows that except for error, observations are consistent (with large eliminations, term constructors can also be observed).

3.1 Status

I have a prototype that has support for all the features of the base language. The fragment for dependent types and type universes satisfies all the conditions above. I believe that all conditions can be made to hold with data types as well.

3.2 Prior work

It is unsurprising that dynamic equality shares many of the same concerns as the large amount of work for contracts, hybrid types, gradual types, and blame. In fact, this work could be seen as gradualizing the Reflection Rule in Extensional Type Theory.

Blame has been strongly advocated for in [29, 28]. Blame tracking can help establish the reasonableness of monitoring systems by linking a dynamic failure directly to the broken static invariant. blame is also a key ingredient of good error messages. However, as many authors have noticed, proving blame correctness is tedious and error prone, it is often only conjectured.

The basic correctness conditions are inspired by the Gradual Guarantee [24]. The implementation also takes inspiration from “Abstracting gradual typing”[11], where static evidence annotations become runtime checks. Unlike some impressive attempts to gradualize the polymorphic lambda calculus [3], dynamic equality does not attempt to preserve any parametric properties of the base language. It is unclear how useful such a restriction to parametric properties would be in practice.

A direct attempt has been made to gradualize a full spectrum dependently typed language to an untyped lambda calculus using the AGT philosophy in [10]. However that system retains the intentional style of equality and user defined data types are not supported. The paper is largely concerned with establishing decidable type checking via an approximate term normalization.

A refinement type system with higher order features is gradualized in [31] though it does not appear powerful enough to be characterized as a full-spectrum dependent type theory. [31] builds on earlier refinement type system work, which described itself as “dynamic”. A notable example is [21] which describes a refinement system that limit’s predicates to base types.

4 Future Work

Dynamic equality appears to be a prerequisite to other interesting possibilities

4.1 Prototyping proofs and programs

Just as “obvious” equalities are missing from the definitional relation, “obvious” proofs and programs are not always conveniently available to the programmer. For instance, in Agda it is possible to write a sorting function quickly using simple types. With effort it is possible to prove that sorting procedure correct by rewriting it with the necessarily invariants. However very little is offered in between. The problem is magnified if module boundaries hide the implementation details of a function, since the details are exactly what is needed to make a proof! This is especially important for larger scale software where a library may require proof terms that while “correct” are not constructable from the exports of other libraries.

The solution proposed here is some additional syntax that will search for a term of the type when resolved at runtime. Given the sorting function

$$\text{sort} : \text{List } \mathbb{N} \rightarrow \text{List } \mathbb{N}$$

and given the first order predicate that

$$\text{IsSorted} : \text{List } \mathbb{N} \rightarrow *$$

then it is possible to assert that `sort` behaves as expected with

$$\lambda x. ? : (x : \text{List } \mathbb{N}) \rightarrow \text{IsSorted}(\text{sort } x)$$

this term will act like any other term at runtime, given a list input it will verify that the `sort` function correctly handles that input, give an error, or non-terminate.

Additionally this would allow simple prototyping from first order specification. For instance,

$$\begin{aligned} & \text{data Mult} : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N} \rightarrow * \text{ where} \\ & \text{base} : (x : \mathbb{N}) \rightarrow \text{Mult} 0 x 0 \\ & \text{suc} : (x y z : \mathbb{N}) \rightarrow \text{Mult} x y z \rightarrow \text{Mult} (1 + x) y (y + z) \end{aligned}$$

can be used to prototype

$$\text{div} = \lambda x. \lambda y. \text{fst} \left(? : \sum z : \mathbb{N}. \text{Mult} x y z \right)$$

The term search can be surprisingly subtle, for instance

$$? : \sum f : \mathbb{N} \rightarrow \mathbb{N}. \text{Id} (f, \lambda x. x + 1) \& \text{Id} (f, \lambda x. 1 + x)$$

depends on the definitional properties of functions. To avoid this subtly I plan to only support term search over first order data.

Though the proof search is currently primitive, better search methods could be incorporated in future work.

4.1.1 Prior work

Proof search is often used for static term generation in dependently typed languages (for instance Coq tactics). A first order theorem prover is attached to Agda in [20].

Twelf made use of runtime proof search but the underling theory cannot be considered full-spectrum.

4.2 Testing dependent programs

Both dynamic equalities and dynamic proof search vastly weaken the guarantees of normal dependent type systems. Programmers still would like a evidence of correctness, even while they intend to provide full proofs of properties in the future. However, there are few options available in full-spectrum dependently typed languages aside from costly and sometimes unconstructable proofs.

The mainstream software industry has similar needs for evidence of correctness, and has made use of testing done in a separate execution phase. Given the rich specifications that dependent types provide it is possible to improve on the hand crafted tests used by most of the industry. Instead we can use a type directed symbolic execution, to run questionable equalities over concrete values and engage and precompute the searched proof terms. Precomputed proof terms can be cached, so that exploration is not too inefficient in the common case of repeating tests at regular intervals of code that is mostly the same. Precomputed terms can be made available at runtime, covering for the inefficient search procedure.

Interestingly dynamic equality is necessary for testing like this, since otherwise, definitional properties of functions would need to be accounted for. Using dynamic equality it is possible only consider the extensional behavior of functions.

Finally future work can add more advanced methods of testing and proof generation. This architecture should make it easier to add more advanced exploration and search without changing the underlining definitional behavior.

4.2.1 Prior work

4.2.2 Symbolic Execution

Most research for Symbolic Execution targets popular languages (like C) and uses SMT solvers to efficiently explore conditional branches that depend on base types. Most work does not support higher order functions or makes simplifying assumptions about the type system. There are however some relevant papers:

- [12] presents a symbolic execution engine supporting Haskell's lazy execution and type system. Higher order functions are not handled
- The draft work[30], handles higher order functions as and inputs provides a proof of completeness
- Symbolic execution for higher order functions for a limited untyped variant of PCF is described in [19]

4.2.3 Testing dependent types

There has been a long recognized need for testing in addition to proving in dependent type systems

- In [9] a QuickCheck style framework was added to an earlier version of Agda
- QuickChick⁵ [8][17, 16, 15] is a research project to add testing to Coq. However testing requires building types classes that establish the properties needed by the testing framework such as decidable equality. This is presumably out of reach of novice Coq users.

References

- [1] Danel Ahman. Fibred computational effects. *arXiv preprint arXiv:1710.02594*, 2017.
- [2] Danel Ahman. Handling fibred algebraic effects. *Proceedings of the ACM on Programming Languages*, 2(POPL):1–29, 2017.
- [3] Amal Ahmed, Dustin Jamner, Jeremy G. Siek, and Philip Wadler. Theorems for free for free: Parametricity, with and without types. *Proc. ACM Program. Lang.*, 1(ICFP), August 2017.
- [4] Lennart Augustsson. Cayenne a language with dependent types. In *Proceedings of the Third ACM SIGPLAN International Conference on Functional Programming*, ICFP ’98, pages 239–250, New York, NY, USA, 1998. Association for Computing Machinery.
- [5] Chris Casinghino. Combining proofs and programs. 2014.
- [6] Chris Casinghino, Vilhelm Sjöberg, and Stephanie Weirich. Combining proofs and programs in a dependently typed language. *ACM SIGPLAN Notices*, 49(1):33–45, 2014.
- [7] Jesper Cockx, Nicolas Tabareau, and Théo Winterhalter. The taming of the rew: A type theory with computational assumptions. *Proceedings of the ACM on Programming Languages*, 2021.
- [8] Maxime Dénès, Catalin Hritcu, Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C Pierce. Quickchick: Property-based testing for coq. In *The Coq Workshop*, 2014.
- [9] Peter Dybjer, Qiao Haiyan, and Makoto Takeyama. Combining testing and proving in dependent type theory. In *International Conference on Theorem Proving in Higher Order Logics*, pages 188–203. Springer, 2003.

⁵<https://github.com/QuickChick/QuickChick>

- [10] Joseph Eremondi, Éric Tanter, and Ronald Garcia. Approximate normalization for gradual dependent types. *Proc. ACM Program. Lang.*, 3(ICFP), July 2019.
- [11] Ronald Garcia, Alison M. Clark, and Éric Tanter. Abstracting gradual typing. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '16, pages 429–442, New York, NY, USA, 2016. Association for Computing Machinery.
- [12] William T. Hallahan, Anton Xue, Maxwell Troy Bland, Ranjit Jhala, and Ruzica Piskac. Lazy counterfactual symbolic execution. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI 2019, pages 411–424, New York, NY, USA, 2019. Association for Computing Machinery.
- [13] Limin Jia, Jianzhou Zhao, Vilhelm Sjöberg, and Stephanie Weirich. Dependent types and program equivalence. In *Proceedings of the 37th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 275–286, 2010.
- [14] Garrin Kimmell, Aaron Stump, Harley D Eades III, Peng Fu, Tim Sheard, Stephanie Weirich, Chris Casinghino, Vilhelm Sjöberg, Nathan Collins, and Ki Yung Ahn. Equational reasoning about programs with general recursion and call-by-value semantics. In *Proceedings of the sixth workshop on Programming languages meets program verification*, pages 15–26, 2012.
- [15] Leonidas Lampropoulos. Random testing for language design. 2018.
- [16] Leonidas Lampropoulos, Diane Gallois-Wong, Cătălin Hrițcu, John Hughes, Benjamin C Pierce, and Li-yao Xia. Beginner’s luck: a language for property-based generators. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*, pages 114–129, 2017.
- [17] Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C Pierce. Generating good generators for inductive relations. *Proceedings of the ACM on Programming Languages*, 2(POPL):1–30, 2017.
- [18] Per Martin-Löf. An intuitionistic theory of types. In Giovanni Sambin and Jan M. Smith, editors, *Twenty-five years of constructive type theory (Venice, 1995)*, volume 36 of *Oxford Logic Guides*, pages 127–172. Oxford University Press, 1998.
- [19] Phuc C Nguyen, Sam Tobin-Hochstadt, and David Van Horn. Higher order symbolic execution for contract verification and refutation. *Journal of Functional Programming*, 27, 2017.

- [20] Ulf Norell. *Towards a practical programming language based on dependent type theory*, volume 32. Citeseer, 2007.
- [21] Xinming Ou, Gang Tan, Yitzhak Mandelbaum, and David Walker. Dynamic typing with dependent types. In Jean-Jacques Levy, Ernst W. Mayr, and John C. Mitchell, editors, *Exploring New Frontiers of Theoretical Informatics*, pages 437–450, Boston, MA, 2004. Springer US.
- [22] Pierre-Marie Pédro and Nicolas Tabareau. The fire triangle how to mix substitution, dependent elimination, and effects. 2020.
- [23] Eric L. Seidel, Ranjit Jhala, and Westley Weimer. Dynamic witnesses for static type errors (or, ill-typed programs usually go wrong). In *Proceedings of the 21st ACM SIGPLAN International Conference on Functional Programming*, ICFP 2016, pages 228–242, New York, NY, USA, 2016. Association for Computing Machinery.
- [24] Jeremy G. Siek, Michael M. Vitousek, Matteo Cimini, and John Tang Boyland. Refined Criteria for Gradual Typing. In Thomas Ball, Rastislav Bodik, Shriram Krishnamurthi, Benjamin S. Lerner, and Greg Morrisett, editors, *1st Summit on Advances in Programming Languages (SNAPL 2015)*, volume 32 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 274–293, Dagstuhl, Germany, 2015. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.
- [25] Vilhelm Sjöberg. A dependently typed language with nontermination. 2015.
- [26] Vilhelm Sjöberg, Chris Casinghino, Ki Yung Ahn, Nathan Collins, Harley D Eades III, Peng Fu, Garrin Kimmell, Tim Sheard, Aaron Stump, and Stephanie Weirich. Irrelevance, heterogeneous equality, and call-by-value dependent type systems. *Mathematically Structured Functional Programming*, 76:112–162, 2012.
- [27] Vilhelm Sjöberg and Stephanie Weirich. Programming up to congruence. In *Proceedings of the 42nd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 369–382, 2015.
- [28] Philip Wadler. A Complement to Blame. In Thomas Ball, Rastislav Bodik, Shriram Krishnamurthi, Benjamin S. Lerner, and Greg Morrisett, editors, *1st Summit on Advances in Programming Languages (SNAPL 2015)*, volume 32 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 309–320, Dagstuhl, Germany, 2015. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.

- [29] Philip Wadler and Robert Bruce Findler. Well-typed programs can't be blamed. In Giuseppe Castagna, editor, *Programming Languages and Systems*, pages 1–16, Berlin, Heidelberg, 2009. Springer Berlin Heidelberg.
- [30] Shu-Hung You, Robert Bruce Findler, and Christos Dimoulas. Dynamic symbolic execution of higher-order functions, 2020.
- [31] Jakub Zalewski, James McKinna, J. Garrett Morris, and Philip Wadler. $\hat{\text{I}}\gg\text{db}$: Blame tracking at higher fidelity. 2020. First ACM SIGPLAN Workshop on Gradual Typing 2020, WGT 2020 ; Conference date: 19-01-2020 Through 25-01-2020.